

# 3

## Using the SQL Trace Facility and TKPROF

### Objectives

After completing this lesson, you should be able to:

- Configure the SQL Trace facility to collect session statistics
- Discuss steps needed before tracing
- Enable application tracing
- Consolidate SQL trace files by using the `trcsess` utility
- Format trace files by using the `tkprof` utility
- Interpret the output of the `tkprof` command
- Verify the SQL problem by using a TKProf Report

## Using the SQL Trace Facility: Overview

- Identify the most impacted session.
- Enable tracing on the identified session.
- Run the application coordinating with the user.
- Measure the client's response time for the operation.
- Disable tracing.
- Gather the trace file from the "DIAGNOSTIC\_DEST" location.
- Generate a TKProf Report and sort the SQL statements in order of most elapsed time by using the following command:

```
OS> tkprof <trace file name> <output file name>  
sort=fchela,exeela,prsela
```

## Steps Needed Before Tracing

- Determine the location for diagnostic traces.
- Choose the most important affected sessions:
  - Find sessions with the highest CPU consumption.
  - Find sessions with the highest waits of a certain type.
  - Find sessions with the highest DB time (10g or later).
- Choose the most important affected clients, services, modules, actions, users, or sessions through Enterprise Manager (if possible) or through user feedback.

## Location for Diagnostic Traces

### DIAGNOSTIC\_DEST

Diagnostic Data	Previous Location	ADR Location
Foreground process traces	USER_DUMP_DEST	\$ADR_HOME/trace
Background process traces	BACKGROUND_DUMP_DEST	\$ADR_HOME/trace
Alert log data	BACKGROUND_DUMP_DEST	\$ADR_HOME/alert \$ADR_HOME/trace
Core dumps	CORE_DUMP_DEST	\$ADR_HOME/cdump
Incident dumps	USER_DUMP_DEST BACKGROUND_DUMP_DEST	\$ADR_HOME/incident/incdir_n

### V\$DIAG\_INFO

**\$ADR\_HOME/trace** <= Oracle Database 11g trace – critical error trace

## Highest CPU Consumption: Example

Find sessions with the highest CPU consumption.

```
SELECT s.sid, s.serial#, p.spid as "OS PID", s.username,
       s.module, st.value/100 as "CPU sec"
FROM v$sesstat st, v$statname sn, v$session s, v$process p
WHERE sn.name = 'CPU used by this session'
AND st.statistic# = sn.statistic#
AND st.sid = s.sid
AND s.paddr = p.paddr
AND s.last_call_et < 1800
AND s.logon_time > (SYSDATE - 240/1440)
ORDER BY st.value;
```

## Highest Waits of a Certain Type: Example

Find sessions with the highest waits of a certain type.

```
SELECT s.sid, s.serial#, p.spid as "OS PID", s.username,
       s.module, se.time_waited
FROM v$session_event se, v$session s, v$process p
WHERE se.event = '&event_name'
AND s.last_call_et < 1800
AND s.logon_time > (SYSDATE - 240/1440)
AND se.sid = s.sid
AND s.paddr = p.addr
ORDER BY se.time_waited;

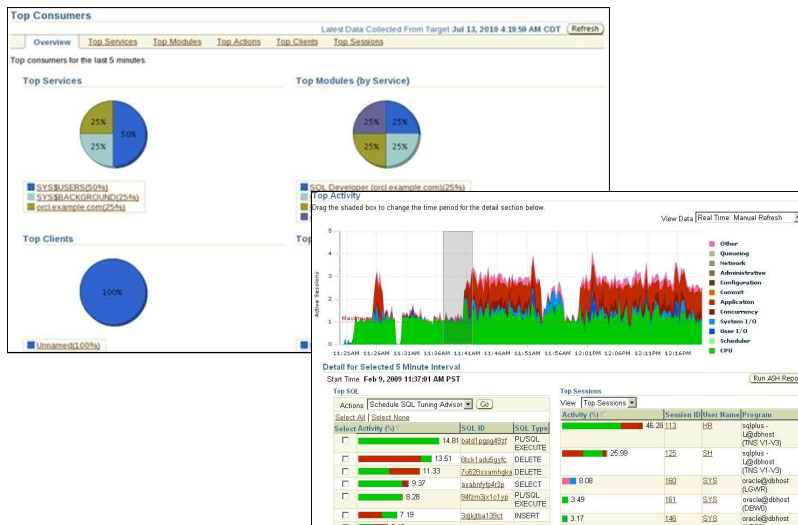
Enter value for event_name: db file sequential read
```

## Highest DB Time: Example

Find sessions with the highest DB time (10g or later).

```
SELECT s.sid, s.serial#, p.spid as "OS PID", s.username,
       s.module, st.value/100 as "DB Time (sec)",
       stcpu.value/100 as "CPU Time (sec)",
       round(stcpu.value / st.value * 100,2) as "%CPU"
FROM v$sesstat st, v$statname sn, v$session s,
     v$sesstat stcpu, v$statname sncpu, v$process p
WHERE sn.name = 'DB time'
AND st.statistic# = sn.statistic#
AND st.sid = s.sid
AND sncpu.name = 'CPU used by this session'
AND stcpu.statistic# = sncpu.statistic#
AND stcpu.sid = st.sid
AND s.paddr = p.addr
AND s.last_call_et < 1800
AND s.logon_time > (SYSDATE - 240/1440)
AND st.value > 0;
```

## Enterprise Manager: Example



## Available Tracing Tools: Overview

- SQL Trace at session level
- ORADEBUG at session level
- LOGON Trigger at a specific user level
- Enterprise Manager
- DBMS\_MONITOR
- DBMS\_APPLICATION\_INFO
- DBMS\_SERVICE
- DBMS\_SESSION
- SQLTXPLAIN (MOS: 215187.1)

## Trace Your Own Session with SQL

Set an event—10046 optimizer trace.

```
SQL> ALTER SESSION SET EVENTS  
2> '10046 trace name context forever, level 12';
```

Execute the statement of interest.

```
SQL> select *  
2> from hr.employees natural join hr.departments  
3> where department_id = 10;
```

Find and view the trace file.

**When to Use:** *It can be used where the session is accessible to the user prior to the start of the statements to be traced.*

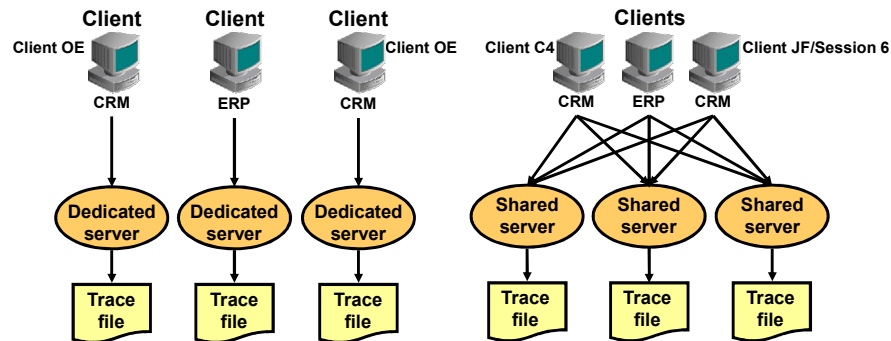
## Trace with a Logon Trigger

**When to Use:** *There may be some situations where it is necessary to trace the activity of a specific user.*

```
CREATE OR REPLACE TRIGGER SYS.set_trace  
AFTER LOGON ON DATABASE  
WHEN (USER like '&USERNAME')  
DECLARE  
    lcommand varchar(200);  
BEGIN  
EXECUTE IMMEDIATE 'alter session set statistics_level=ALL';  
EXECUTE IMMEDIATE 'alter session set max_dump_file_size=UNLIMITED';  
DBMS_MONITOR.SESSION_TRACE_ENABLE(waits=> true, binds=> true);  
END set_trace;  
/
```

Find and view the trace file.

## Consideration: Tracing Challenge



- I want to retrieve traces from CRM service.
- I want to retrieve traces from client C4.
- I want to retrieve traces from session 6.

## What Is a Service?

- Is a means of grouping sessions that perform the same kind of work
- Provides a single-system image instead of a multiple-instances image
- Is a part of the regular administration tasks that provide dynamic service-to-instance allocation
- Is the base for high availability of connections
- Provides a performance-tuning dimension
- Is a handle for capturing trace information

## Using Services with Client Applications

```
ERP= (DESCRIPTION=
      (ADDRESS= (PROTOCOL=TCP) (HOST=mynode) (PORT=1521))
      (CONNECT_DATA= (SERVICE_NAME=ERP)))
```

```
url="jdbc:oracle:oci:@ERP"
```

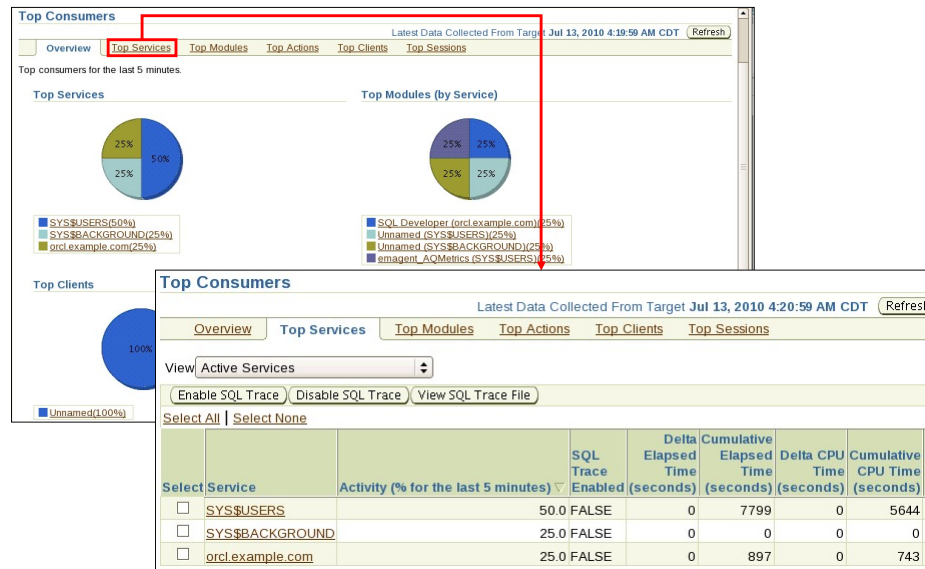
```
url="jdbc:oracle:thin:@ (DESCRIPTION=
      (ADDRESS= (PROTOCOL=TCP) (HOST=mynode) (PORT=1521))
      (CONNECT_DATA= (SERVICE_NAME=ERP))) "
```

## End-to-End Application Tracing

- Simplifies the process of diagnosing performance problems in multitier environments by allowing application workloads to be seen by:
  - Service
  - Module
  - Action
  - Session
  - Client
- End-to-end application tracing tools:
  - Enterprise Manager
  - DBMS\_APPLICATION\_INFO, DBMS\_SERVICE, DBMS\_MONITOR, DBMS\_SESSION
  - SQL Trace and trcsess utility
  - tkprof



## Trace with Enterprise Manager



## Trace with DBMS\_MONITOR

- Automatic service aggregation level of statistics
- DBMS\_MONITOR used for finer granularity of service aggregations:
  - SERV\_MOD\_ACT\_STAT\_ENABLE
  - SERV\_MOD\_ACT\_STAT\_DISABLE
- Possible additional aggregation levels:
  - SERVICE\_NAME/MODULE
  - SERVICE\_NAME/MODULE/ACTION
- Tracing services, modules, and actions:
  - SERV\_MOD\_ACT\_TRACE\_ENABLE
  - SERV\_MOD\_ACT\_TRACE\_DISABLE
- Database settings persist across instance restarts.

## Service Tracing: Example

- Trace on service, module, and action:

```
exec DBMS_MONITOR.SERV_MOD_ACT_TRACE_ENABLE('AP');
```

```
exec DBMS_MONITOR.SERV_MOD_ACT_TRACE_ENABLE(-  
  'AP', 'PAYMENTS', 'QUERY_DELINQUENT');
```

- Trace a particular client identifier:

```
exec DBMS_MONITOR.CLIENT_ID_TRACE_ENABLE  
  (client_id=>'C4', waits => TRUE, binds => FALSE);
```

## Session Tracing: Example

- For all sessions in the database:

```
EXEC dbms_monitor.DATABASE_TRACE_ENABLE(TRUE,TRUE);
```

```
EXEC dbms_monitor.DATABASE_TRACE_DISABLE();
```

- For a particular session:

```
EXEC dbms_monitor.SESSION_TRACE_ENABLE(session_id=>  
27, serial_num=>60, waits=>TRUE, binds=>FALSE);
```

```
EXEC dbms_monitor.SESSION_TRACE_DISABLE(session_id  
=>27, serial_num=>60);
```

## Trace Your Own Session

- Enabling trace:

```
EXEC DBMS_SESSION.SESSION_TRACE_ENABLE(waits =>
TRUE, binds => FALSE);
```

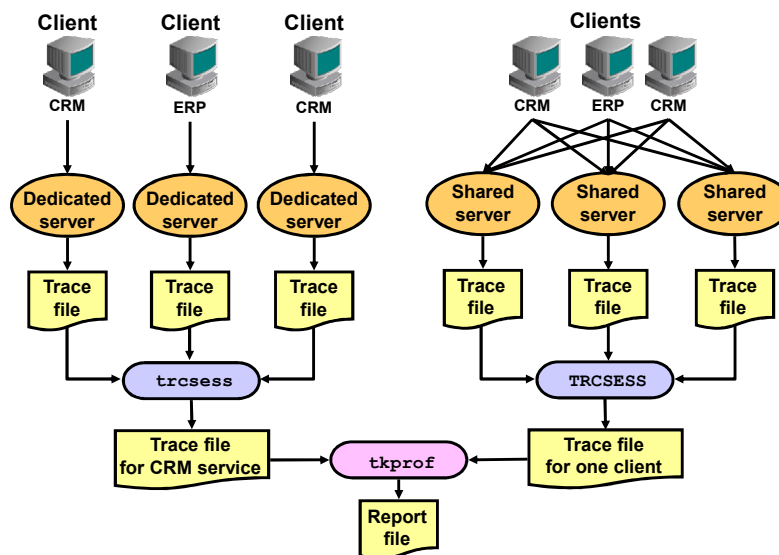
- Disabling trace:

```
EXEC DBMS_SESSION.SESSION_TRACE_DISABLE();
```

- Easily identifying your trace files:


```
alter session set
tracefile_identifier='mytraceid';
```

## The trcsess Utility



## Invoking the trcsess Utility

```
trcsess [output=output_file_name]
       [session=session_id]
       [clientid=client_identifier]
       [service=service_name]
       [action=action_name]
       [module=module_name]
       [<trace file names>]
```



## The trcsess Utility: Example

```
exec dbms_session.set_identifier('HR session');
```

First session

Second session

```
exec dbms_session.set_identifier('HR session');
```

Third session

```
exec Dclient_id=>'HR session', waits => FALSE, -
binds => FALSE);
DBMS_MONITOR.CLIENT_ID_TRACE_ENABLE( -
```

```
select * from employees;
```

```
select * from departments;
```

...

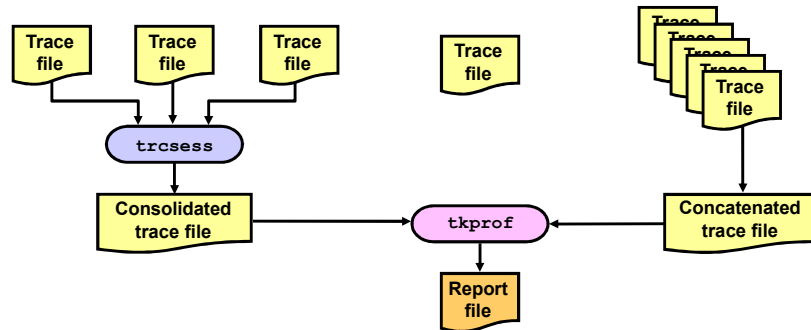
```
exec DBMS_MONITOR.CLIENT_ID_TRACE_DISABLE( -
client_id => 'HR session');
```

```
trcsess output=mytrace.trc clientid='HR session'
$ORACLE_BASE/diag/rdbms/orcl/orcl/trace/*.trc
```

## Formatting SQL Trace Files: Overview


Use the `tkprof` utility to format your SQL trace files:

- Sort raw trace file to exhibit top SQL statements.
- Filter dictionary statements.



## Invoking the tkprof Utility

```
tkprof inputfile outputfile [waits=yes|no]
                                [sort=option]
                                [print=n]
                                [aggregate=yes|no]
                                [insert=sqlscriptfile]
                                [sys=yes|no]
                                [table=schema.table]
                                [explain=user/password]
                                [record=statementfile]
                                [width=n]
```



### tkprof Sorting Options

Sort Option	Description
prscnt	Number of times parse was called
prscpu	CPU time parsing
prsela	Elapsed time parsing
prsdsk	Number of disk reads during parse
prsqry	Number of buffers for consistent read during parse
prscu	Number of buffers for current read during parse
prsmis	Number of misses in the library cache during parse
execnt	Number of executes that were called
execpu	CPU time spent executing
exeela	Elapsed time executing
exedsk	Number of disk reads during execute
exeqry	Number of buffers for consistent read during execute
execu	Number of buffers for current read during execute

## tkprof Sorting Options

Sort Option	Description
exerow	Number of rows processed during execute
exemis	Number of library cache misses during execute
fchcnt	Number of times fetch was called
fchcpu	CPU time spent fetching
fchela	Elapsed time fetching
fchdsk	Number of disk reads during fetch
fchqry	Number of buffers for consistent read during fetch
fchcu	Number of buffers for current read during fetch
fchrow	Number of rows fetched
userid	User ID of user that parsed the cursor

## TKProf Report Structure

```

...
select max(cust_credit_limit) from customers where cust_city = 'Paris'

```

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.02	0.02	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	2	0.00	0.00	0	15	0	1
total	4	0.02	0.02	0	15	0	1


```

Misses in library cache during parse: 1
Optimizer mode: FIRST_ROWS
Parsing user id: 88

Rows      Row Source Operation
-----
1  TABLE ACCESS FULL EMPLOYEES (cr=15 r=0 w=0 time=1743 us)
1  SORT AGGREGATE (cr=7 r=0 w=0 time=777 us)
107 TABLE ACCESS FULL EMPLOYEES (cr=7 r=0 w=0 time=655 us)

Elapsed times include waiting on following events:
Event waited on                      Times Max. Wait Total Waited
-----
SQL*Net message to client              2      0.00      0.00
SQL*Net message from client            2      9.62      9.62

```





## Interpret a TKProf Report: Example

- Row source plan

Rows	Row Source Operation
-----	
[A] 1	[B] TABLE ACCESS FULL EMPLOYEES ([C] cr=15 [D] r=0 [E] w=0 [F] time=1743 us)
1	SORT AGGREGATE (cr=7 r=0 w=0 time=777 us)
107	TABLE ACCESS FULL EMPLOYEES (cr=7 r=0 w=0 time=655 us)

- Spotting relatively high resource usage

update ...							
where ...							
-----							
call	count	cpu	elapsed	disk	query	current	rows
-----							
Parse	1	7	122	0	0	0	0
Execute	1	75	461	5	[H] 297	[I] 3	[J] 1
Fetch	0	0	0	0	0	0	1
-----							
total	2	82	583	5	297	3	2

## Interpret a TKProf Report: Example

Spotting overparsing

```
select ...
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	[M] 2	[N] 221	329	0	45	0	0
Execute	3	[O] 9	[P] 17	0	0	0	0
Fetch	3	6	8	0	[L] 4	0	[K] 1
total	8	236	354	0	49	0	1

Misses in library cache during parse: 1 [Q]  
...

## Interpret a TKProf Report: Example

Spotting queries that execute too much

```
update ...  
set ...  
where ...
```

call	count	cpu	elapsed	disk	query	current	rows
Parse	0	0	0	0	0	0	0
Execute	488719	66476.95	66557.80	1	488729	1970566	488719
Fetch	0	0	0	0	0	0	0
total	488719	66476.95	66557.80	1	488729	1970566	488719

...

## What to Verify: Example


Was total elapsed time in TKProf account for the application response time measured when the application was executed?

OVERALL TOTALS FOR ALL NON-RECURSIVE STATEMENTS							
call	count	cpu	elapsed	disk	query	current	rows
Parse	1165	0.66	2.15	0	45	0	0
Execute	2926	1.23	2.92	0	0	0	0
Fetch	2945	117.03	398.23	5548	1699259	16	39654
total	7036	118.92	403.31	5548	1699304	16	39654
OVERALL TOTALS FOR ALL RECURSIVE STATEMENTS							
call	count	cpu	elapsed	disk	query	current	rows
Parse	0	0.00	0.00	0	0	0	0
Execute	0	0.00	0.00	0	0	0	0
Fetch	0	0.00	0.00	0	0	0	0
total	0	0	0.00	0	0	0	0

## What to Verify: Example

Has the time spent parsing, executing, and fetching account for most of the elapsed time recorded in the trace?

call	count	cpu	elapsed	disk	query	current	rows
Parse	1	0.00	0.00	0	0	0	0
Execute	1	0.00	0.00	0	0	0	0
Fetch	8	0.00	0.00	0	14	0	14
total	10	0	0.00	0	14	0	14
...							
Rows	Row Source Operation						
-----							
14	TABLE ACCESS FULL EMPLOYEES (cr=14 r=0 w=0 time=377 us)						
Elapsed times include waiting on following events:							
Event waited on		Times Max. Wait		Total Waited			
		Waited					
SQL*Net message to client		8		0.00		0.00	
SQL*Net message from client		8		16.36		78.39	



### What to Verify: Example

- Is the query you expect to tune shown at the top of the TKProf Report?
- Does the query spend most of its time in the Execute and Fetch phases (not Parse phase)? Make sure the trace file contains data only from the recent test.

call	count	cpu	elapsed	disk	query	current	rows
Parse	555	100.09	300.83	0	0	0	0
Execute	555	0.42	0.78	0	0	0	0
Fetch	555	14.04	85.03	513	1448514	0	11724
total	1665	114.55	368.65	513	1448514	0	11724



### Quiz

In an environment with an applications server that uses a connection pool, you use \_\_\_\_\_ to identify which trace files need to be combined to get an overall trace of the application.

- a. trcsess
- b. tkprof
- c. Oracle SQL Developer
- d. DBMS\_APPLICATION\_INFO

## Summary

In this lesson, you should have learned how to:

- Configure the SQL Trace facility to collect session statistics
- Discuss steps needed before tracing
- Enable application tracing
- Consolidate SQL trace files by using the `trcsess` utility
- Format trace files by using the `tkprof` utility
- Interpret the output of the `tkprof` command
- Verify the SQL problem by using a TKProf Report

## Practice 3: Overview

This practice covers the following topics:

- Creating a service
- Tracing your application by using services
- Interpreting trace information by using `trcsess` and `tkprof`