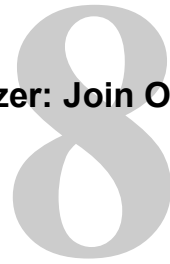


## Optimizer: Join Operators



### Objectives

After completing this lesson, you should be able to:

- Describe the SQL operators for joins
- List the possible access paths

## Join Methods

A join:

- Defines the relationship between two row sources
- Is a method of combining data from two data sources
- Is controlled by join predicates, which define how the objects are related
- Join methods:
  - Nested loops
  - Sort-merge join

```
SELECT e.ename, d.dname
FROM dept d JOIN emp e USING (deptno)
WHERE e.job = 'ANALYST' OR e.empno = 9999;
```

← Join predicate

← Nonjoin predicate

```
SELECT e.ename, d.dname
FROM emp e, dept d
WHERE e.deptno = d.deptno AND
      (e.job = 'ANALYST' OR e.empno = 9999);
```

← Join predicate

← Nonjoin predicate

## Nested Loops Join

- Driving row source is scanned.
- Each row returned drives a lookup in inner row source.
- Joining rows are then returned.



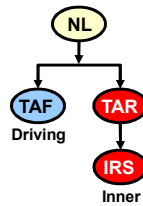
```
select ename, e.deptno, d.deptno, d.dname
from emp e, dept d
where e.deptno = d.deptno and ename like 'A%';
```

Id	Operation	Name	Rows	Cost
0	SELECT STATEMENT		2	4
1	NESTED LOOPS		2	4
2	TABLE ACCESS FULL	EMP	2	2
3	TABLE ACCESS BY INDEX ROWID	DEPT	1	1
4	INDEX UNIQUE SCAN	PK_DEPT	1	1

2 - filter("E"."ENAME" LIKE 'A%')

4 - access("E"."DEPTNO"="D"."DEPTNO")

## Nested Loops Join: Prefetching



```

select ename, e.deptno, d.deptno, d.dname
from emp e, dept d
where e.deptno = d.deptno and ename like 'A%';
  
```

0	SELECT STATEMENT			2	84	5
1	TABLE ACCESS BY INDEX ROWID	DEPT		1	22	1
2	NESTED LOOPS			2	84	5
* 3	TABLE ACCESS FULL	EMP		2	40	3
* 4	INDEX RANGE SCAN	IDEPT		1		0

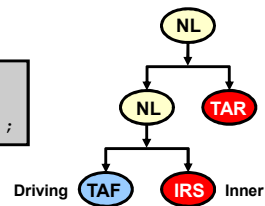
```

3 - filter("E"."ENAME" LIKE 'A%')
4 - access("E"."DEPTNO"="D"."DEPTNO")
  
```

## Nested Loops Join: 11g Implementation

```

select ename, e.deptno, d.deptno, d.dname
from emp e, dept d
where e.deptno = d.deptno and ename like 'A%';
  
```



Explain Plan x

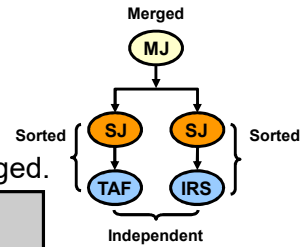
0.299 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			4	1
NESTED LOOPS				
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates				
ENAME LIKE 'A%'				
INDEX	PK_DEPT	UNIQUE SCAN	0	1
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	DEPT	BY INDEX RO...	1	1

## Sort-Merge Join

- First and second row sources are sorted by the same sort key.
- Sorted rows from both tables are merged.

```
select /*+ USE_MERGE(d e) NO_INDEX(d) */
  ename, e.deptno, d.deptno, dname
from emp e, dept d
where e.deptno = d.deptno and ename > 'A'
```

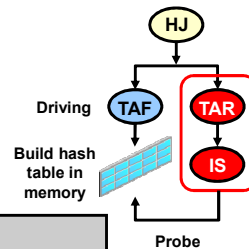


OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY...
SELECT STATEMENT			8	14
MERGE JOIN			8	14
SORT			4	4
TABLE ACCESS	DEPT	FULL	3	4
SORT			4	14
Access Predicates				
AND				
E.DEPTNO=D.DEPTNO				
SYS_OP_DESCEND(E.DEPT)				
Filter Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	EMP	FULL	3	14
Filter Predicates				
ENAME>'A'				

## Hash Join

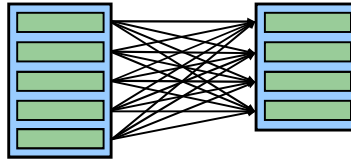
- The smallest row source is used to build a hash table.
- The second row source is hashed and checked against the hash table.

```
select /*+ USE_HASH(e d) */
  ename, e.deptno, d.deptno, dname
from emp e, dept d
where e.deptno = d.deptno and ename like 'A%'
```



OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY...
SELECT STATEMENT			7	1
HASH JOIN			7	1
Access Predicates				
AND				
E.DEPTNO=D.DEPTNO				
SYS_OP_DESCEND(E.DEPTNO)				
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates				
ENAME LIKE 'A%'				
TABLE ACCESS	DEPT	FULL	3	4

## Cartesian Join



```
select ename, e.deptno, d.deptno, dname
from emp e, dept d where ename like 'A%';
```

Explain Plan x

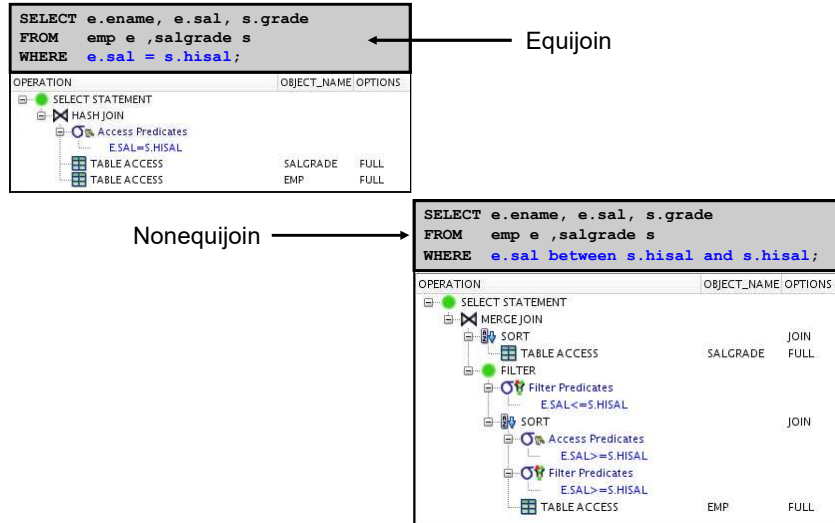
0.295 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			6	4
MERGE JOIN		CARTESIAN	6	4
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates				
ENAME LIKE 'A%'				
BUFFER		SORT	3	4
TABLE ACCESS	DEPT	FULL	3	4

## Join Types

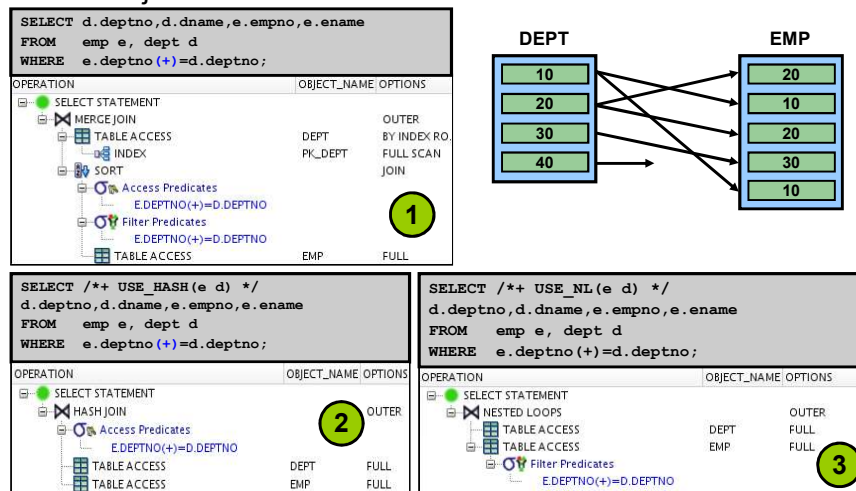
- A join operation combines the output from two row sources and returns one resulting row source.
- Join operation types include the following :
  - Join (Equijoin/Natural – Nonequijoin)
  - Outer join (Full, Left, and Right)
  - Semi join: `EXISTS` subquery
  - Anti join: `NOT IN` subquery
  - Star join (Optimization)

## Equijoins and Nonequijoins



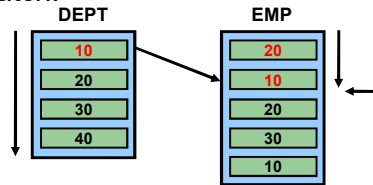
## Outer Joins

An outer join returns a row even if no match is found.

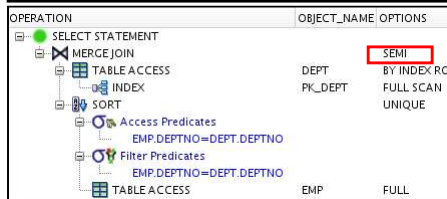


## Semijoins

Semijoins look only for the first match.

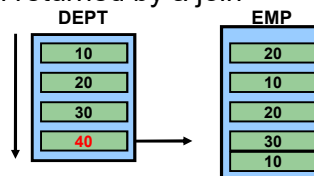


```
SELECT deptno, dname
FROM dept
WHERE EXISTS (SELECT 1 FROM emp WHERE emp.deptno=dept.deptno);
```

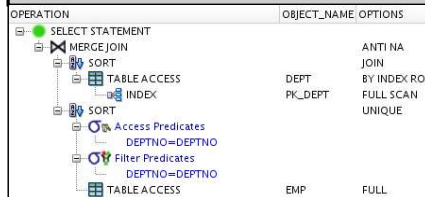


## Antijoins

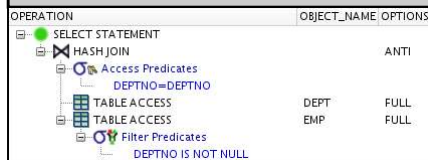
Reverse of what would have been returned by a join



```
SELECT deptno, dname
FROM dept
WHERE deptno not in
(SELECT deptno FROM emp);
```



```
SELECT deptno, dname FROM dept
WHERE deptno IS NOT NULL AND
deptno NOT IN
(SELECT /*+ HASH_AJ */ deptno FROM
emp WHERE deptno IS NOT NULL);
```



### Quiz

The \_\_\_\_\_ join is used when one or more of the tables do not have any join conditions to any other tables in the statement.

- a. Hash
- b. Cartesian
- c. Nonequijoin
- d. Outer

### Quiz

The \_\_\_\_\_ join returns a row even if no match is found.

- a. Hash
- b. Cartesian
- c. Semi
- d. Outer



### Quiz

The \_\_\_\_\_ join looks only for the first match.

- a. Hash
- b. Cartesian
- c. Semi
- d. Outer

### Quiz

In a hash join, the \_\_\_\_\_ row source is used to build a hash table.

- a. Biggest
- b. Smallest
- c. Sorted
- d. Unsorted

## Summary

In this lesson, you should have learned to:

- Describe the SQL operators for joins
- List the possible access paths

## Practice 8: Overview

This practice covers using different join paths for better optimization.