

10

Optimizer Statistics

Objectives

After completing this lesson, you should be able to:

- Describe optimizer statistics
 - Table statistics
 - Index statistics
 - Column statistics (histogram)
 - Column statistics (extended statistics)
 - System statistics
- Gather optimizer statistics
- Set statistic preferences
- Use dynamic sampling
- Manage optimizer statistics
- Discuss optimizer statistics best practices

Optimizer Statistics

- Describe the database and the objects in the database
- Information used by the query optimizer to estimate:
 - Selectivity of predicates
 - Cost of each execution plan
 - Access method, join order, and join method
 - CPU and I/O costs
- Types of optimizer statistics:
 - Table statistics
 - Index statistics
 - Column statistics
 - System statistics

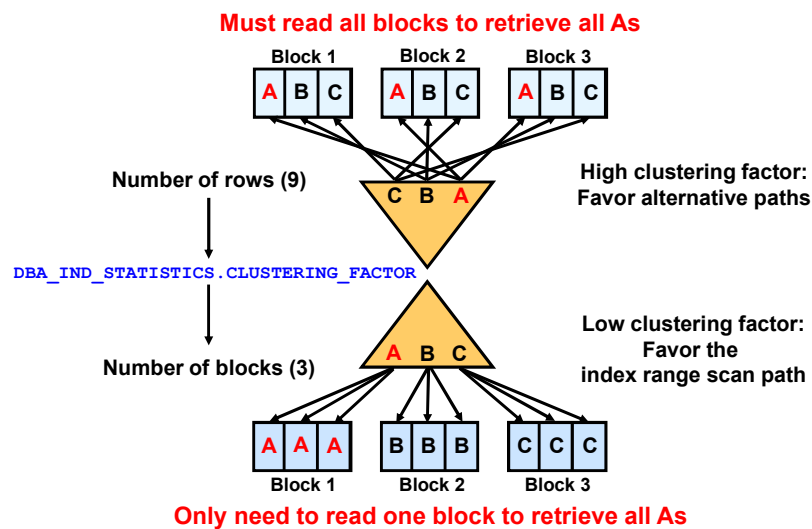
Table Statistics (DBA_TAB_STATISTICS)

- Used to determine:
 - Table access cost
 - Join cardinality
 - Join order
- Some of the table statistics gathered are:
 - Row count (NUM_ROWS)
 - Block count (BLOCKS) *Exact*
 - Average row length (AVG_ROW_LEN)
 - Statistics status (STALE_STATS)

Index Statistics (DBA_IND_STATISTICS)

- Used to decide:
 - Full table scan versus index scan
- Statistics gathered are:
 - B*-tree level (BLEVEL) *Exact*
 - Leaf block count (LEAF_BLOCKS)
 - Clustering factor (CLUSTERING_FACTOR)
 - Distinct keys (DISTINCT_KEYS)
 - Average number of leaf blocks in which each distinct value in the index appears (AVG_LEAF_BLOCKS_PER_KEY)
 - Average number of data blocks in the table pointed to by a distinct value in the index (AVG_DATA_BLOCKS_PER_KEY)
 - Number of rows in the index (NUM_ROWS)

Index Clustering Factor



Column Statistics (DBA_TAB_COL_STATISTICS)


- Count of distinct values of the column (NUM_DISTINCT)
- Low value (LOW_VALUE) *Exact*
- High value (HIGH_VALUE) *Exact*
- Number of nulls (NUM_NULLS)
- Selectivity estimate for nonpopular values (DENSITY)
- Number of histogram buckets (NUM_BUCKETS)
- Type of histogram (HISTOGRAM)

Column Statistics: Histograms

- The optimizer assumes uniform distributions; this may lead to suboptimal access plans in the case of data skew.
- Histograms:
 - Store additional column distribution information
 - Give better selectivity estimates in the case of nonuniform distributions
- With unlimited resources, you could store each different value and the number of rows for that value.
- This becomes unmanageable for a large number of distinct values, and a different approach is used:
 - Frequency histogram (#distinct values ≤ #buckets)
 - Height-balanced histogram (#buckets < #distinct values)
- They are stored in DBA_TAB_HISTOGRAMS.

Frequency Histograms

10 buckets, 10 distinct values



Distinct values: 1, 3, 5, 7, 10, 16, 27, 32, 39, 49

Number of rows: 40001

Viewing Frequency Histograms

```
BEGIN
  DBMS_STATS.gather_table_STATS (OWNNAME=>'OE', TABNAME=>'INVENTORIES',
    METHOD_OPT => 'FOR COLUMNS SIZE 20 warehouse_id');
END;
```

```
SELECT column_name, num_distinct, num_buckets, histogram
FROM   USER_TAB_COL_STATISTICS
WHERE  table_name = 'INVENTORIES' AND
       column_name = 'WAREHOUSE_ID';
```


COLUMN_NAME	NUM_DISTINCT	NUM_BUCKETS	HISTOGRAM
WAREHOUSE_ID	9	9	FREQUENCY

```
SELECT endpoint_number, endpoint_value
FROM   USER_HISTOGRAMS
WHERE  table_name = 'INVENTORIES' and column_name = 'WAREHOUSE_ID'
ORDER BY endpoint_number;
```

ENDPOINT_NUMBER	ENDPOINT_VALUE
36	1
213	2
261	3
...	...

Height-Balanced Histograms

5 buckets, 10 distinct values
(8000 rows per bucket)



Distinct values: 1, 3, 5, 7, 10, 16, 27, 32, 39, 49

Number of rows: 40001

Viewing Height-Balanced Histograms

```
BEGIN
  DBMS_STATS.gather_table_stats(ownname => 'OE', tabname => 'INVENTORIES',
    method_opt => 'FOR COLUMNS SIZE 10 quantity_on_hand');
END;
```

```
SELECT column_name, num_distinct, num_buckets, histogram
FROM USER_TAB_COL_STATISTICS
WHERE table_name = 'INVENTORIES' AND column_name = 'QUANTITY_ON_HAND';
```

COLUMN_NAME	NUM_DISTINCT	NUM_BUCKETS	HISTOGRAM
QUANTITY_ON_HAND	237	10	HEIGHT BALANCED

```
SELECT endpoint_number, endpoint_value
FROM USER_HISTOGRAMS
WHERE table_name = 'INVENTORIES' and column_name = 'QUANTITY_ON_HAND'
ORDER BY endpoint_number;
```

ENDPOINT_NUMBER	ENDPOINT_VALUE
0	0
1	27
2	42
3	57
...	

Best Practices: Histogram

- Histograms are useful when you have a high degree of skew in the column distribution.
- Histograms are *not* useful for:
 - Columns which do not appear in the `WHERE` or `JOIN` clauses
 - Columns with uniform distributions
 - Equality predicates with unique columns
- The maximum number of buckets is the least (254, # distinct values). If possible, frequency histograms are preferred.
- Do not use histograms unless they substantially improve performance.


Best Practices: Histogram

- Set `METHOD_OPT` to `FOR ALL COLUMNS AUTO`.
- Use `TRUNCATE` instead of dropping and re-creating the same table if you need to remove all rows from a table.
- When upgrading to 11g, use the same histograms used initially in earlier releases.
- If incorrect cardinality / selectivity is observed in an execution plan, check to see if a histogram can resolve the problem.
- Make sure statistics for objects are collected at the highest sample size you can afford and see if the plan improves.
- In earlier releases, if a query uses binds or binds are not representative of future executions, we should not consider histograms to avoid bind peeking. In 11g, adaptive cursor sharing resolves bind/histogram issues.

Column Statistics: Extended Statistics

- The optimizer poorly estimates selectivity *on Highly Correlated Column Predicates*:
 - Columns have values that are highly correlated.
 - Actual selectivity is often much lower or higher than the optimizer estimates. For example,
`WHERE cust_state_province = 'CA'`
`AND country_id=52775;`
- The optimizer poorly estimates *Expression on Columns*:
 - `WHERE upper(model)='MODEL'`
 - When a function is applied to a column in the `WHERE` clause, the optimizer has no way of knowing how that function affects the selectivity of the column.
- In these cases, a group of columns within a table or an expression on a column can be gathered to obtain a more accurate selectivity value.

Multicolumn Statistics




```
select
dbms_stats.create_extended_stats('jfv','vehicle','(make,model)')
from dual;
```

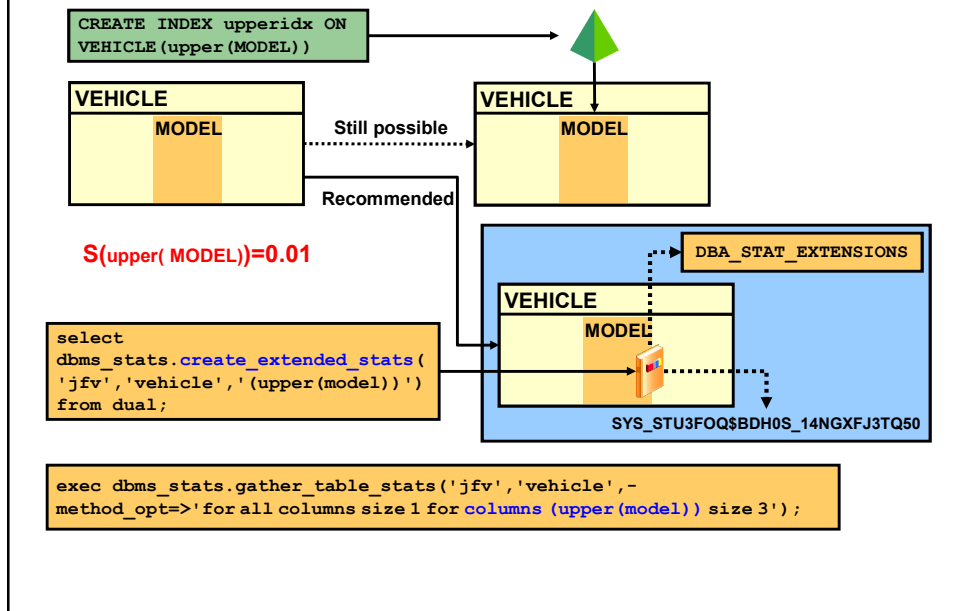
2

```
exec dbms_stats.gather_table_stats('jfv','vehicle',-
method_opt=>'for all columns size 1 for columns (make,model) size 3');
```

3



Expression Statistics



System Statistics

- System statistics are used to estimate:
 - I/O performance and utilization
 - CPU performance and utilization
- System statistics enable the query optimizer to estimate I/O and CPU costs more accurately, enabling the query optimizer to choose a better execution plan.
- Procedures:
 - DBMS_STATS.GATHER_SYSTEM_STATS
 - DBMS_STATS.SET_SYSTEM_STATS
 - DBMS_STATS.GET_SYSTEM_STATS

System Statistics: Example

Viewing System Statistics:

```
SELECT * FROM sys.aux_stats$;
```

SNAME	PNAME	PVAL1	PVAL2
SYSSTATS_INFO	STATUS		COMPLETED
SYSSTATS_INFO	DSTART		08-09-2001 16:40
SYSSTATS_INFO	DSTOP		08-09-2001 16:42
SYSSTATS_INFO	FLAGS	0	
SYSSTATS_MAIN	SREADTIM	7.581	
SYSSTATS_MAIN	MREADTIM	56.842	
SYSSTATS_MAIN	CPUSPEED	117	
SYSSTATS_MAIN	MBRC	9	

Best Practices: System Statistics

- System statistics must be gathered on a regular basis; this does not invalidate cached plans.
- Gathering system statistics equals analyzing system activity for a specified period of time.
- When gathering the optimizer system statistics:
 - It is highly recommended that you gather system statistics during normal workload for several hours.
 - If no real workload is available, you can also gather `NORWORKLOAD` statistics.

Gathering System Statistics: Automatic Collection Example

First day

```
EXECUTE DBMS_STATS.GATHER_SYSTEM_STATS(  
  interval => 120,  
  stattab => 'mystats', statid => 'OLTP');
```

First night

```
EXECUTE DBMS_STATS.GATHER_SYSTEM_STATS(  
  interval => 120,  
  stattab => 'mystats', statid => 'OLAP');
```

Next days

```
EXECUTE DBMS_STATS.IMPORT_SYSTEM_STATS(  
  stattab => 'mystats', statid => 'OLTP');
```

Next nights

```
EXECUTE DBMS_STATS.IMPORT_SYSTEM_STATS(  
  stattab => 'mystats', statid => 'OLAP');
```

Gathering System Statistics: Manual Collection Example

- Start manual system statistics collection in the data dictionary:

```
EXECUTE DBMS_STATS.GATHER_SYSTEM_STATS( -  
  gathering_mode => 'START');
```

- Generate the workload.
- End the collection of system statistics:

```
EXECUTE DBMS_STATS.GATHER_SYSTEM_STATS( -  
  gathering_mode => 'STOP');
```

Gathering Statistics: Overview

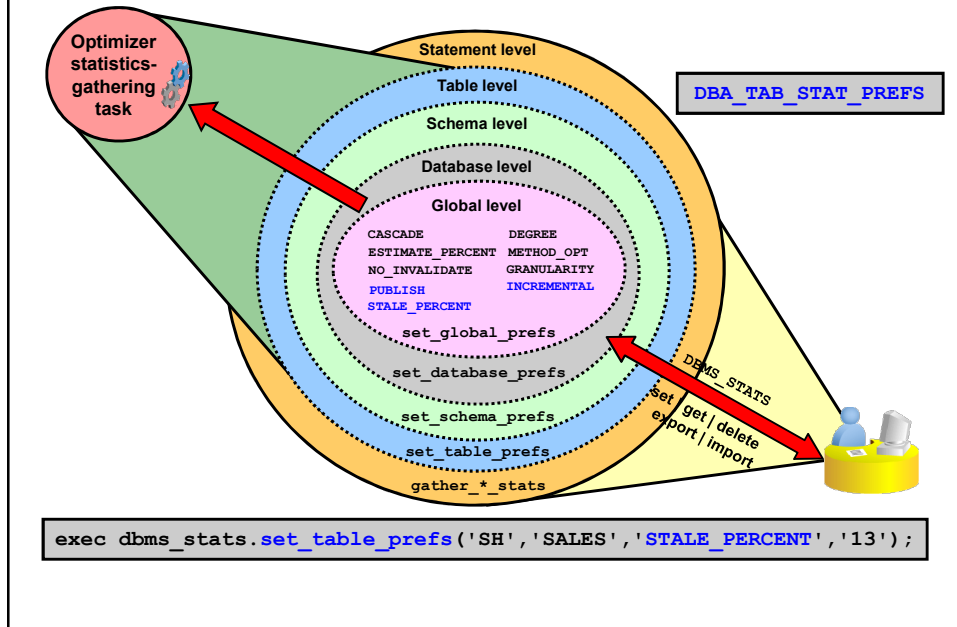
- Automatic statistics gathering
 - `gather_stats_prog` automated task
- Manual statistics gathering
 - `DBMS_STATS` package
- Dynamic sampling
- When statistics are missing
- Vendor-recommended gathering

Selectivity:	
Equality	1%
Inequality	5%
Other predicates	5%
Table row length	20
# of index leaf blocks	25
# of distinct values	100
Table cardinality	100
Remote table cardinality	2000

Automatic Statistics Gathering

- Oracle's recommended method for collecting statistics
- Oracle Database 11g automates optimizer statistics collection:
 - Statistics are gathered automatically only on all database objects that have no statistics or have stale statistics (> 10% of rows modified)
 - The `gather_stats_prog` automated task is used for statistics collection and maintenance.
- Automated statistics collection:
 - Eliminates need for manual statistics collection
 - Significantly reduces the chances of poor execution plans
- The Statistic Preferences feature is available in Oracle 11g for some objects that require statistics collection settings that are different from the database default.

Statistic Preferences: Overview



Manual Statistics Gathering

You can use Enterprise Manager and the `DBMS_STATS` package to:

- Generate and manage statistics for use by the optimizer:
 - Gather/Modify
 - View/Name
 - Export/Import
 - Delete/Lock
- Gather statistics on:
 - Indexes, tables, columns, partitions
 - Object, schema, or database
- Gather statistics either serially or in parallel
- gather/set system statistics (currently not possible in EM)

When to Gather Statistics Manually

- Rely mostly on automatic statistics collection:
 - Change the frequency of automatic statistics collection to meet your needs.
 - Remember that `STATISTICS_LEVEL` should be set to `TYPICAL` or `ALL` for automatic statistics collection to work properly.
- Gather statistics manually for:
 - Objects that are volatile
 - Objects modified in batch operations (Gather statistics as part of the batch operation.)
 - External tables, system statistics, fixed objects
 - New objects (Gather statistics right after object creation.)

Manual Statistics Collection: Factors

- Monitor objects for DMLs.
- Determine the correct sample sizes.
- Determine the degree of parallelism.
- Determine if histograms should be used.
- Determine the cascading effects on indexes.
- Procedures to use in `DBMS_STATS`:
 - `GATHER_INDEX_STATS`
 - `GATHER_TABLE_STATS`
 - `GATHER_SCHEMA_STATS`
 - `GATHER_DICTIONARY_STATS`
 - `GATHER_DATABASE_STATS`
 - `GATHER_SYSTEM_STATS`

Gathering Object Statistics: Example

```
dbms_stats.gather_table_stats
('sh'           -- schema
,'customers'    -- table
, null         -- partition
, 20           -- sample size(%)
, false        -- block sample?
,'for all columns' -- column spec
, 4            -- degree of parallelism
,'default'     -- granularity
, true );      -- cascade to indexes
```

```
dbms_stats.set_param('CASCADE',
                    'DBMS_STATS.AUTO_CASCADE');
dbms_stats.set_param('ESTIMATE_PERCENT', '5');
dbms_stats.set_param('DEGREE', 'NULL');
```

Best Practices: Object Statistics

- Ensure that all objects (tables and indexes) have statistics gathered.
- Use a sample size that is large enough if feasible.
- Gather optimizer statistics during periods of low activity.
- If partitions are in use, gather global statistics if possible.
- Use Oracle Database 11g pending statistics to verify effect of new statistics when tuning to minimize risk.
- Gather statistics after data has been loaded (>10% added), but before indexes are created.

Optimizer Dynamic Sampling: Overview

- Dynamic sampling can be done for tables and indexes:
 - Without statistics
 - Whose statistics cannot be trusted, starting with 11gR2 if object statistics are stale and sampling level => 4
- Used to determine more accurate statistics when estimating:
 - Table cardinality
 - Predicate selectivity
- Feature controlled by:
 - `OPTIMIZER_DYNAMIC_SAMPLING` parameter
 - `OPTIMIZER_FEATURES_ENABLE` parameter
 - `DYNAMIC_SAMPLING` hint
 - `DYNAMIC_SAMPLING_EST_CDN` hint

Optimizer Dynamic Sampling at Work

- Sampling is done at compile time.
- If a query benefits from dynamic sampling:
 - A recursive SQL statement is executed to sample data.
 - The number of blocks sampled depends on the `OPTIMIZER_DYNAMIC_SAMPLING` initialization parameter.
- During dynamic sampling, predicates are applied to the sample to determine selectivity.
- Use dynamic sampling when:
 - Sampling time is a small fraction of the execution time (like Data Warehouse, not OLTP).
 - Volatile data is used with `DELETE_*_STATS` and `LOCK_*_STATS`.
 - Correlated columns are used in the `WHERE` clause.
 - Global temporary tables are used.
 - The query is executed many times.
 - You believe a better plan can be found (during testing).

OPTIMIZER_DYNAMIC_SAMPLING

- Dynamic session or system parameter.
- Can be set to a value from "0" to "10."
- "0" turns off dynamic sampling.
- "1" samples all unanalyzed tables, if an unanalyzed table:
 - Is joined to another table or appears in a subquery or nonmergeable view
 - Has no indexes
 - Has more than 32 blocks
- "2" samples all unanalyzed tables.
- The higher the value, the more aggressive application of sampling.
- Dynamic sampling is repeatable if no update activity occurred.


Managing Statistics: Overview (Export / Import / Lock / Restore / Publish)

- Purpose:
 - To revert to preanalyzed statistics if gathering statistics causes critical statements to perform badly
 - To test the new statistics before publishing
- Importing previously exported statistics (9i)
- Locking statistics on a specific table (10g)
- Restoring statistics archived before gathering (10g)
- Statistics can be pending before publishing (11gR2)

Export and Import Statistics

Use DBMS_STATS procedures:

- CREATE_STAT_TABLE creates the statistics table.
- EXPORT_*_STATS moves the statistics to the statistics table.
- Use Data Pump to move the statistics table.
- IMPORT_*_STATS moves the statistics to the data dictionary.



Locking Statistics

- Prevents automatic gathering
- Is mainly used for volatile tables:
 - Lock without statistics implies dynamic sampling.

```
BEGIN
  DBMS_STATS.DELETE_TABLE_STATS('OE','ORDERS');
  DBMS_STATS.LOCK_TABLE_STATS('OE','ORDERS');
END;
```

- Lock with statistics for representative values.

```
BEGIN
  DBMS_STATS.GATHER_TABLE_STATS('OE','ORDERS');
  DBMS_STATS.LOCK_TABLE_STATS('OE','ORDERS');
END;
```

- The FORCE argument overrides statistics locking.

```
SELECT stattype_locked FROM dba_tab_statistics;
```

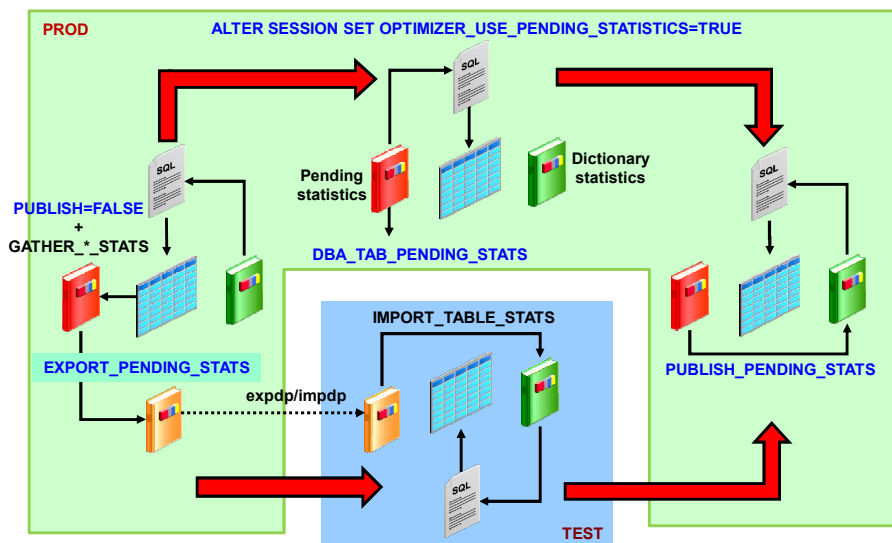
Restoring Statistics

- Past statistics may be restored with the DBMS_STATS.RESTORE_*_STATS procedures.

```
BEGIN
  DBMS_STATS.RESTORE_TABLE_STATS(
    OWNNAME=>'OE', TABNAME=>'INVENTORIES',
    AS_OF_TIMESTAMP=>'15-JUL-10 09.28.01.597526000 AM -05:00');
END;
```

- Statistics are automatically stored:
 - With the timestamp in DBA_TAB_STATS_HISTORY
 - When collected with DBMS_STATS procedures
- Statistics are purged:
 - When the STATISTICS_LEVEL is set to TYPICAL or ALL automatically
 - After 31 days or time defined by DBMS_STATS.ALTER_STATS_HISTORY_RETENTION

Deferred Statistics Publishing: Overview



Deferred Statistics Publishing: Example

```
exec dbms_stats.set_table_prefs('SH','CUSTOMERS','PUBLISH','false'); ①
```

```
exec dbms_stats.gather_table_stats('SH','CUSTOMERS'); ②
```

```
alter session set optimizer_use_pending_statistics = true; ③
```

Execute your workload from the same session. ④

```
exec dbms_stats.publish_pending_stats('SH','CUSTOMERS'); ⑤
```

Quiz

When there are no statistics for an object being used in a SQL statement, the optimizer uses:

- a. Rule-based optimization
- b. Dynamic sampling
- c. Fixed values
- d. Statistics gathered during the parse phase
- e. Random values

Quiz

The optimizer depends on accurate statistics to produce the best execution plans. The automatic statistics-gathering task does not gather statistics on everything. Which objects require you to gather statistics manually?

- a. External tables
- b. Data dictionary
- c. Fixed objects
- d. Volatile tables
- e. System statistics

Quiz

There is a very volatile table in the database. The size of the table changes by more than 50 percent daily. What steps are part of the procedure to force dynamic sampling?

- a. Delete statistics.
- b. Lock statistics.
- c. Gather statistics when the table is at its largest.
- d. Set `DYNAMIC_SAMPLING=9`.
- e. Set `DYNAMIC_SAMPLING=0`.
- f. Allow the `DYNAMIC_SAMPLING` parameter to default.

Summary

In this lesson, you should have learned how to:

- Describe optimizer statistics
 - Table statistics
 - Index statistics
 - Column statistics (histogram)
 - Column statistics (extended statistics)
 - System statistics
- Gather optimizer statistics
- Set statistic preferences
- Use dynamic sampling
- Manage optimizer statistics
- Discuss optimizer statistics best practices

Practice 10: Overview

This practice covers the following topics:

- Using system statistics
- Using automatic statistics gathering